

Parallel Gauss-Seidel method for solving the heat transfer problem

1.2 Approach

The use of iterative methods to solve computational problems is motivated by the need to solve large scale problems such as systems of linear equations with a considerable number of variables. In this cases, the use of a direct method, while providing an exact solution, would require a considerable amount of computing power and time to execute, making iterative methods more suitable for such scenarios. One example of such method is the Gauss-Seidel iterative method that works by continuously updating a matrix containing the coefficients of each equation in each line where each element is updated as follows:

$$x_i^{k+1} = \frac{b_i - \sum_{j=1}^{i-1} a_{ij}x_j^{k+1} - \sum_{j=i+1}^n a_{ij}x_j^k}{a_{ii}}$$

where k is the current iteration.

The method halts when the differences between iterations are under a predefined threshold/tolerance. This improves on the Jacobi method by using the values of the elements that have already been computed to update the current element, resulting in a higher convergence rate as well dispensing the need for storing two vectors as only the most recent values are needed. However, by using the most recent values, this method introduces data dependencies in each iteration, making it harder to develop in a parallel paradigm (as we will see in a later section, this can be overcome by applying a strategy known as Red-Black strategy).

Another important factor to take into consideration when solving large scale problems is the use of matrix-free methods. Such approaches don't require the (coefficient) matrix to be explicitly stored in memory, as that would require a large amount of memory and computing time, accessing the coefficients by using matrix-vector products. However, for the sake of simplicity, the implementations developed in the present work use a matrix to store the coefficients of each equation in each line of the matrix.

1.3 Discretization

Calculating the temperature at every point in a 2D plane in a continuous domain would be infeasible as it would require an infinite amount of points. As such, it becomes necessary to transfer the domain from a continuous one to a discrete one through a process known as discretization. This process works by applying a mesh to the continuous domain and taking the value of the function (being studied) at each point of the mesh. In the case of the heat equation the mesh is composed of four boundaries, called heat sources, with known temperature values, and a set of internal points whose temperature is unknown. The process of discretization introduces a truncation error due to the fact that less points than the initial problem are used to describe the 2D plane. As such, the truncation error is inversely proportional to the amount of points in the mesh.




Figure 1: Discretization on a 2D plane with a five-point stencil

This truncation error in a mathematical perspective it's a swap of partial derivatives for stencils (approximation), i.e. use of finite differences. As such, each point will be updated according to a five-point stencil, using the arithmetic average of the four neighbouring points and the central point as the new value. In other words, with a centered finite differences approach the truncation error will be $O(h^2)$ where h is the distance between two consecutive points in the mesh ($h = 1$ when using Gauss-Seidel).

2 Benchmark Platform

In order to better understand the benchmark results gathered a brief description of the conditions in which the benchmarks were executed will be presented.

2.1 Statistical Methods

A sample size of one was used for each benchmark because the number of iterations required don't change between different executions of the same benchmark.

2.2 Hardware Platform

The benchmarks hereby mentioned were executed on a cluster node, from the SeARCH Computing Infrastructure. The node used, r641, presents the following hardware specifications [2, 1]:

- **CPU:** Intel Xeon E5-2650v2 @2.6GHz, 16 cores
- **Cache Size:**
 - L1d (per core): 32 KiB
 - L1i (per core): 32 KiB
 - L2 (per core): 256 KiB
 - L3 (per package): 20 MiB
- **RAM:** 64GB
- **Network:** Ethernet, *myrinet*

3 Sequential Implementation

When implementing problems such as the heat transfer problem that deal with large quantities of data, and operations per element/unit of data, it's important to consider an approach that allows for easy parallelization. However, in order to better understand the problem being solved, a sequential implementation of the problem was developed initially, from which the pseudo-parallel (or parallel-ready) implementation was based off.

The sequential implementation (see A.2) begins by defining a tri-dimensional matrix made up of two N by N matrices that store the temperature values of the plate at each iteration. It begins by initializing the top, left and right boundaries of the plate at 100 and the bottom one with 0 while the internal points are initialized with a value of 50. Based on the size of the plate **i.e.** number of points in the mesh(N), a tolerance is set for the difference between temperatures in consecutive iterations. The function `poissons` implements the Gauss-Seidel method on the matrix and halts the loop when the stopping criteria is met:

$$||u^{k+1} - u^k|| \leq TOL$$

At each iteration, the last temperature values are stored in the matrix indexed as `plate[last]`. Additionally, the biggest difference between consecutive temperatures in a point, is stored in `dif` which is used to check whether the stopping criteria has been met.

4 Sequential Implementation with Red-Black Strategy

Based on the sequential implementation, a parallel-ready implementation was developed (see A.3) using the Red-Black strategy for dealing with data dependencies in each iteration. This strategy works by interleaving the updated points in each loop so that it's possible for the loop load to be distributed amongst different processes without needing to deal with data dependencies. It achieves this by updating the matrix in two different loops, one for updating the odd points(Red) and the other for updating the even points(Black). Even though there are no dependencies in each

of the loops, the order at which the loops are executed must be preserved (Red-Black). A *caveat* to this mechanism is the fact that the Red points stencil uses the old values for temperature:

$$x_i^{k+1} = \frac{b_i - \sum_{j=1}^n a_{ij} x_j^k}{a_{ii}}$$

while the Black points stencil uses the current temperature values of it's neighboring elements:

$$x_i^{k+1} = \frac{b_i - \sum_{j=1}^n a_{ij} x_j^{k+1}}{a_{ii}}$$

This property results in a bigger number of iterations needed as the method tends to converge at a slower rate. The stop criteria is the same as the previously stated.

5 Parallel Implementation (MPI) with Red-Black Strategy

The development of parallel applications is usually divided into two different paradigms the shared memory paradigm and the distributed memory paradigm. The choice between the two is influenced not only by the characteristics of the work load, but also by the intended result. The development of parallel applications using a distributed memory paradigm via libraries such as OpenMPI allows for a thorough theoretical analysis of the expected behaviour of the algorithm, which is useful for identifying possible bottlenecks before developing the application. The parallel implementation developed uses OpenMPI and distributes the load between each process. Due to the characteristics of OpenMPI, this (initial) distribution doesn't require a broadcast of any values to each process as their memory is copied from the parent process, which contains the matrix with the initial temperatures for each point in the mesh. The load distribution process begins by calculating the number of lines assigned to each process, which is given by the expression: $\frac{N-2}{\text{number_processes}}$ where N is the number of lines. The process with rank $P - 1$ is assigned with the remaining lines **i.e.** $(N - 2) \bmod P$:




Figure 2: Load distribution

Each process then invokes the **phase** function which updates a group of rows of the mesh based on the process rank. To ensure the order at which the mesh points are updated complies with the Red-Black scheme, the function

`updateValues` is called after updating the Red points, and each process exchanges one line with the two neighbouring processes (rank-1 and rank+1):




Figure 3: Function `updateValues`

At the end of each iteration, the function `updateValues` is invoked again and every process communicates the biggest difference between two consecutive temperatures for a given point to the parent process (with rank 0), which then determines the biggest value between all this differences. This value is then communicated back to each process which use it to check whether the method has reached an acceptable solution. After all iterations are executed, it is necessary to collect (communicate) the data from all the processes, which is gathered by the parent process.




Figure 4: Result gathering

The total amount of data exchanged between processes can be estimated as a function of the mesh size N , the number of processes P and the number of iterations it by the following expression:

$$it * (2 * 2N * (P - 1) + 2(P - 1)) + \frac{N - 2}{P} * N * (P - 1)$$

where:

- $2 * 2N * (P - 1)$ corresponds to the amount of data transferred in each iteration
- $2(P - 1)$ corresponds to the amount of data needed to determine the biggest difference between two consecutive temperature values
- $\frac{N-2}{P} * N * (P - 1)$ corresponds to the amount of data gathered after all iterations are executed

From this expression it's obvious that the amount of data transferred is of order $O(N^2)$ based on the mesh size.

6 Analysis and comparison of the different approaches

To understand the behaviour of each implementation, and more specifically, how the Red-Black strategy affects the number of iterations needed to obtain an acceptable solution, each implementation was run with 3 different parameters for the mesh size (10, 100, 1000) as well as different tolerance values: $TOL = \frac{1}{N^2}$. The following chart shows the number of iterations needed, for each implementation, to reach an acceptable solution based on this parameters:




Figure 5: Results

One noticeable aspect of the results is the fact that the Red-Black strategy requires a higher number of iterations to obtain an acceptable solution, which comes as a consequence of the fact that the Red points are updated using old values for it's neighbouring points which are then used by the Black points, resulting in a slower convergence rate. Additionally, an increment on the size of the mesh, N , leads to a higher number of iterations until the stop criteria is met. This is to be expected as the tolerance value is a function of the mesh size, becoming more restrictive as this value increases. Thus, the convergence rate is lower for bigger N values, while the cost of each iteration increases since there are more elements to be updated. The number of iterations increases proportional to a logarithmic scale. Finally, the use of a parallel version via MPI, which distributes the load between more than one process (Parallel with RB), while requiring the same number of iterations as it's sequential counterpart (Sequential with RB), will result in a reduced execution time.

7 Conclusion

The experiments on this work show that the usage of strategy such as Red-Black, in conjunction with an iterative method such as Gauss-Seidel, for solving the heat dispersion problem is indeed advantageous for dealing with data dependencies that would, otherwise, make parallelization infeasible. This remains true even with an increase in the number of iterations needed to reach an acceptable solution as the workload distribution outweighs this increase when it comes to performance and execution time. Additionally, the use of an iterative method such as Gauss-Seidel, allows for an acceptable solution to be obtained for linear systems that would otherwise take a considerable amount of execution time and computing power to solve.

References

- [1] *cluster nodes*. http://search6.di.uminho.pt/wordpress/?page_id=55. Access in 02/05/2019.
- [2] *Intel Xeon E5-2650 v2 - CM8063501375101 / BX80635E52650V2*. <http://www.cpu-world.com/CPUs/Xeon/Intel-Xeon%20E5-2650%20v2.html>. Access in 02/05/2019.

Appendix A Code

A.1 Auxiliar Code (utils)

```
#include "utils.h"

/**
 * Initialize heat plate with temperature values
 */
void initPlate(float plate[][N][N]){
    for(int i = 0; i < N; i++){
        plate[0][0][i] = 100.0f;
        plate[0][i][0] = 100.0f;
        plate[0][i][N-1] = 100.0f;
        plate[0][N-1][i] = 0.0f;

        plate[1][0][i] = 100.0f;
        plate[1][i][0] = 100.0f;
        plate[1][i][N-1] = 100.0f;
        plate[1][N-1][i] = 0.0f;
    }

    for(int i = 1; i < N-1; i++){
        for(int j = 1; j < N-1; j++){
            plate[0][i][j] = 50.0f;
        }

        plate[0][N-1][0] = 0.0f;
        plate[1][N-1][0] = 0.0f;
    }
}

void initPlateForMPI(float plate[][N*N]){
    for(int i = 0; i < N; i++){
        plate[0][0*N+i] = 100.0f;
        plate[0][i*N+0] = 100.0f;
        plate[0][i*N+N-1] = 100.0f;
        plate[0][(N-1)*N+i] = 0.0f;

        plate[1][0*N+i] = 100.0f;
        plate[1][i*N+0] = 100.0f;
        plate[1][i*N+N-1] = 100.0f;
        plate[1][(N-1)*N+i] = 0.0f;
    }

    for(int i = 1; i < N-1; i++){
        for(int j = 1; j < N-1; j++){
            plate[0][i*N+j] = 50.0f;
            plate[1][i*N+j] = 50.0f;
        }

        plate[0][(N-1)*N+0] = 0.0f;
        plate[1][(N-1)*N+0] = 0.0f;
    }
}
```


A.2 Sequential Implementation

```
#include "utils.h"

/*
 * Returns the number of iterations performed
 */
int poissongs(float plate[][N][N], float tol){
    int it = 0, last=0;
    float dif = tol + 1.0f, temp;

    while(dif > tol){
        dif = 0.0f; //differences at the borders are 0.0f
        for(int i = 1; i < N-1; i ++){
            for(int j = 1; j < N-1; j ++){
                plate[!last][i][j] = (plate[!last][i-1][j] + plate[!last][i][j-1]
                                     + plate[last][i][j+1] + plate[last][i+1][j]) / 4.0f;
                temp = fabs(plate[!last][i][j] - plate[last][i][j]);
                if(temp > dif) dif = temp;
            }
            it ++;
            last = !last; //update matrix
        }
        return it;
    }
}

int main(int argc, char *argv[]){
    float plate[2][N][N];
    float tol = 1.0f / (N*N);
    int it;

    initPlate(plate);

    it = poissongs(plate, tol);
    printf("Sequential Poisson GS\n Iteration Count: %d\n", it);

    return 0;
}
```

A.3 Sequential Implementation with Red-Black Strategy

```
#include "utils.h"

/*
 * Returns the number of iterations performed
 */
int poissongs(float plate[][N][N], float tol){
```

```

int it = 0, last=0;
float dif = tol + 1.0f, temp;

while(dif > tol){
    dif = 0.0f;           //differences at the borders are 0.0f
    //Black
    for(int i = 1; i < N-1; i ++){
        for(int j = 2 - i%2; j < N-1; j +=2){
            plate[!last][i][j] = (plate[last][i-1][j] + plate[last][i][j-1]
                                   + plate[last][i][j+1] + plate[last][i+1][j]) / 4.0f;
            temp = fabs(plate[!last][i][j] - plate[last][i][j]);
            if(temp > dif) dif = temp;
        }

        //Red
        for(int i = 1; i < N-1; i ++){
            for(int j = 1 + i%2; j < N-1; j +=2){
                plate[!last][i][j] = (plate[!last][i-1][j] + plate[!last][i][j-1]
                                       + plate[!last][i][j+1] + plate[!last][i+1][j]) / 4.0f;
                temp = fabs(plate[!last][i][j] - plate[last][i][j]);
                if(temp > dif) dif = temp;
            }

            it ++;
            last = !last; //update matrix
        }

        return it;
    }
}

int main(int argc, char *argv[]){
    float plate[2][N][N];
    float tol = 1.0f / (N*N);
    int it;

    initPlate(plate);

    it = poissongs(plate, tol);

    printf("Sequential Poisson GS with Red-Back strategy\n Iteration Count: %d\n", it);

    return 0;
}

```

A.4 Parallel Implementation (MPI) with Red-Black Strategy

```

#include <mpi.h>
#include "utils.h"

```

```

void updateValues(float plate[][N*N], int last, int offset, int rows, int rank, int no_procs, MPI_Status status){
    if(!rank){
        MPI_Recv( &(plate[last][(rows+1)*N]), N, MPI_FLOAT, rank+1, 2, MPI_COMM_WORLD, &status );
        MPI_Send( &(plate[last][rows*N]), N, MPI_FLOAT, rank+1, 2, MPI_COMM_WORLD);
    }else{
        MPI_Send( &(plate[last][(offset+1)*N]), N, MPI_FLOAT, rank-1, 2, MPI_COMM_WORLD);

        if(rank!=no_procs-1)
            MPI_Recv( &(plate[last][offset*N + (rows+1)*N]), N, MPI_FLOAT, rank+1, 2, MPI_COMM_WORLD, &status );

        if(rank!=no_procs-1)
            MPI_Send( &(plate[last][offset*N + rows*N]), N, MPI_FLOAT, rank+1, 2, MPI_COMM_WORLD);

        MPI_Recv( &(plate[last][offset*N]), N, MPI_FLOAT, rank-1, 2, MPI_COMM_WORLD, &status );
    }
}

float phase(float plate[][N*N], int last, int offset, int rows, int rank, int no_procs, MPI_Status status){
    float dif, temp;
    dif = 0.0f;

    //Black
    for(int i = offset+1; i < offset + 1 + rows; i ++){
        for(int j = 2 - i%2; j < N-1; j +=2){
            plate[!last][i*N+j] = (plate[last][(i-1)*N+j] + plate[last][i*N+j-1]
                                   + plate[last][i*N+j+1] + plate[last][(i+1)*N+j]) / 4.0f;
            temp = fabs(plate[!last][i*N+j] - plate[last][i*N+j]);
            if(temp > dif) dif = temp;
        }

        updateValues(plate, !last, offset, rows, rank, no_procs, status);

        //Red
        for(int i = offset+1; i < offset + 1 + rows; i ++){
            for(int j = 1 + i%2; j < N-1; j +=2){
                plate[!last][i*N+j] = (plate[!last][(i-1)*N+j] + plate[!last][i*N+j-1]
                                       + plate[!last][i*N+j+1] + plate[!last][(i+1)*N+j]) / 4.0f;
                temp = fabs(plate[!last][i*N+j] - plate[last][i*N+j]);
                if(temp > dif) dif = temp;
            }

            updateValues(plate, !last, offset, rows, rank, no_procs, status);

            if(rank){
                MPI_Send( &dif, 1, MPI_FLOAT, 0, 1, MPI_COMM_WORLD);
                MPI_Recv( &dif, 1, MPI_FLOAT, 0, 1, MPI_COMM_WORLD, &status);
            }else{
                for(int i=1; i<no_procs; i++){
                    MPI_Recv( &temp, 1, MPI_FLOAT, i, 1, MPI_COMM_WORLD, &status);
                    if(temp > dif) dif = temp;
                }
                for(int i=1; i<no_procs; i++){
                    MPI_Send( &dif, 1, MPI_FLOAT, i, 1, MPI_COMM_WORLD);
                }
            }
        }
        return dif;
    }

    /*
    * Returns the number of iterations performed
    */
    int poissongs(float plate[][N*N], float tol, int rank, int no_procs, MPI_Status status){
        int it = 0, last=0, rows_per_proc, remaining_rows, work_load = N-2, begin, offset;
        float dif = tol + 1.0f;

        rows_per_proc = work_load/(no_procs-1); //truncates result

```

```

remaining_rows = work_load % (no_procs-1);    //remaining rows

offset = rank * rows_per_proc;

while(dif > tol){
    if(rank==no_procs-1)
        dif = phase(plate, last, offset, remaining_rows, rank, no_procs, status);
    else
        dif = phase(plate, last, offset, rows_per_proc, rank, no_procs, status);

    it ++;
    last = !last; //update matrix
}

//collect data
if(!rank){
    begin = (rows_per_proc+1)*N;

    for(int i=1; i<no_procs-1; i++, begin+=rows_per_proc*N)
        MPI_Recv(&(plate[last][begin]), rows_per_proc*N, MPI_FLOAT, i, 3, MPI_COMM_WORLD, &status);

    MPI_Recv(&(plate[last][begin]), remaining_rows*N, MPI_FLOAT, no_procs-1, 3, MPI_COMM_WORLD, &status);

    //copy of process 0 result
    for(int i=1; i<rows_per_proc+1; i++)
        for(int j=0; j<N; j++)
            plate[last][i*N+j]=plate[!last][i*N+j];

}else if(rank == no_procs-1) {
    MPI_Send( &(plate[!last][(offset+1)*N]), remaining_rows*N, MPI_FLOAT, 0, 3, MPI_COMM_WORLD);
}else{
    MPI_Send( &(plate[!last][(offset+1)*N]), rows_per_proc*N, MPI_FLOAT, 0, 3, MPI_COMM_WORLD);
}

return it;
}

int main(int argc, char *argv[]){
    float plate[2][N*N];
    float tol = 1.0f / (N*N);
    int it, rank, no_procs;

    initPlateForMPI(plate);

    //Move MPI process creation to poissongs
    MPI_Status status;
    MPI_Init(&argc, &argv);
    MPI_Comm_size(MPI_COMM_WORLD, &no_procs);
    MPI_Comm_rank( MPI_COMM_WORLD, &rank );

    it = poissongs(plate, tol, rank, no_procs, status);

    MPI_Finalize();

    if(rank == 0)
        printf("Parallel Poisson GS with Red-Back strategy\n Iteration Count: %d\n", it);

    return 0;
}

```